

LoRA: Low-Rank Adaptation of Large Language Models

Edward J. Hu • Yelong Shen • Phillip Wallis • Weizhu Chen

Microsoft Corporation

ABSTRACT

We propose **LoRA (Low-Rank Adaptation)**, which freezes pre-trained model weights and injects trainable rank decomposition matrices into Transformer attention layers, reducing trainable parameters by 10,000x with zero inference latency.

